

# PRNGine: Massively Parallel Pseudo-Random Number Generation and Probability Distribution Approximations on AMD AI Engines

**Mohamed Bouaziz**, Suhaib A. Fahmy

Accelerated Connected Computing Lab (**ACCL**),

King Abdullah University of Science and Technology (**KAUST**), Saudi Arabia

# Pseudo-Random Number Generation/ors

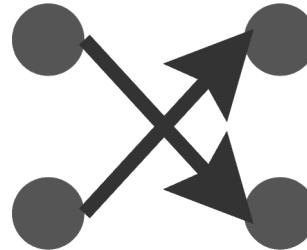
- Random numbers are crucial for running many applications
- Some applications require a massive generation of random numbers in parallel



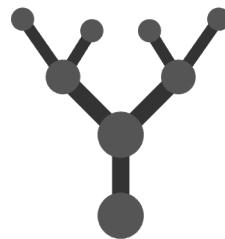
Cryptography



Monte Carlo simulation



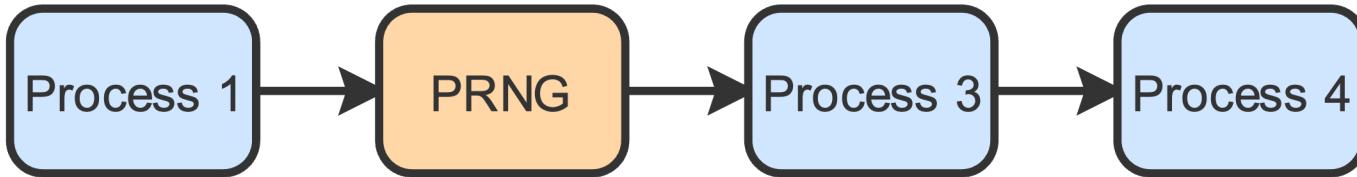
Data shuffling



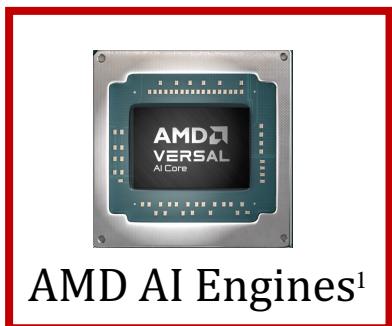
Procedural content generation

# PRNGs in Processing Pipelines

- CGRAs/RDAs are well-suited for running dataflows in pipelines.



- There is an increasing demand for such dataflow accelerators.



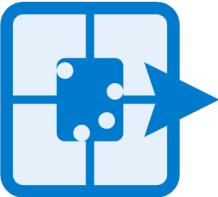
1: <https://www.amd.com/en/products/adaptive-socs-and-fpgas/technologies/ai-engine.html>

2: <https://groq.com/>

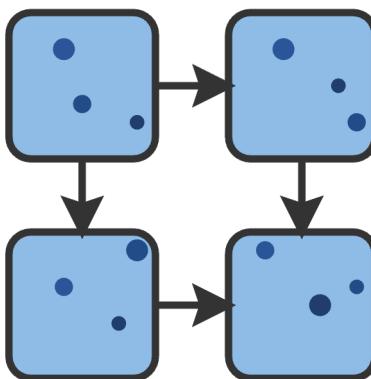
3: <https://sambanova.ai/>

# Existing work

- LUT-optimised generators on FPGAs → fine-grained bitwise manipulations.

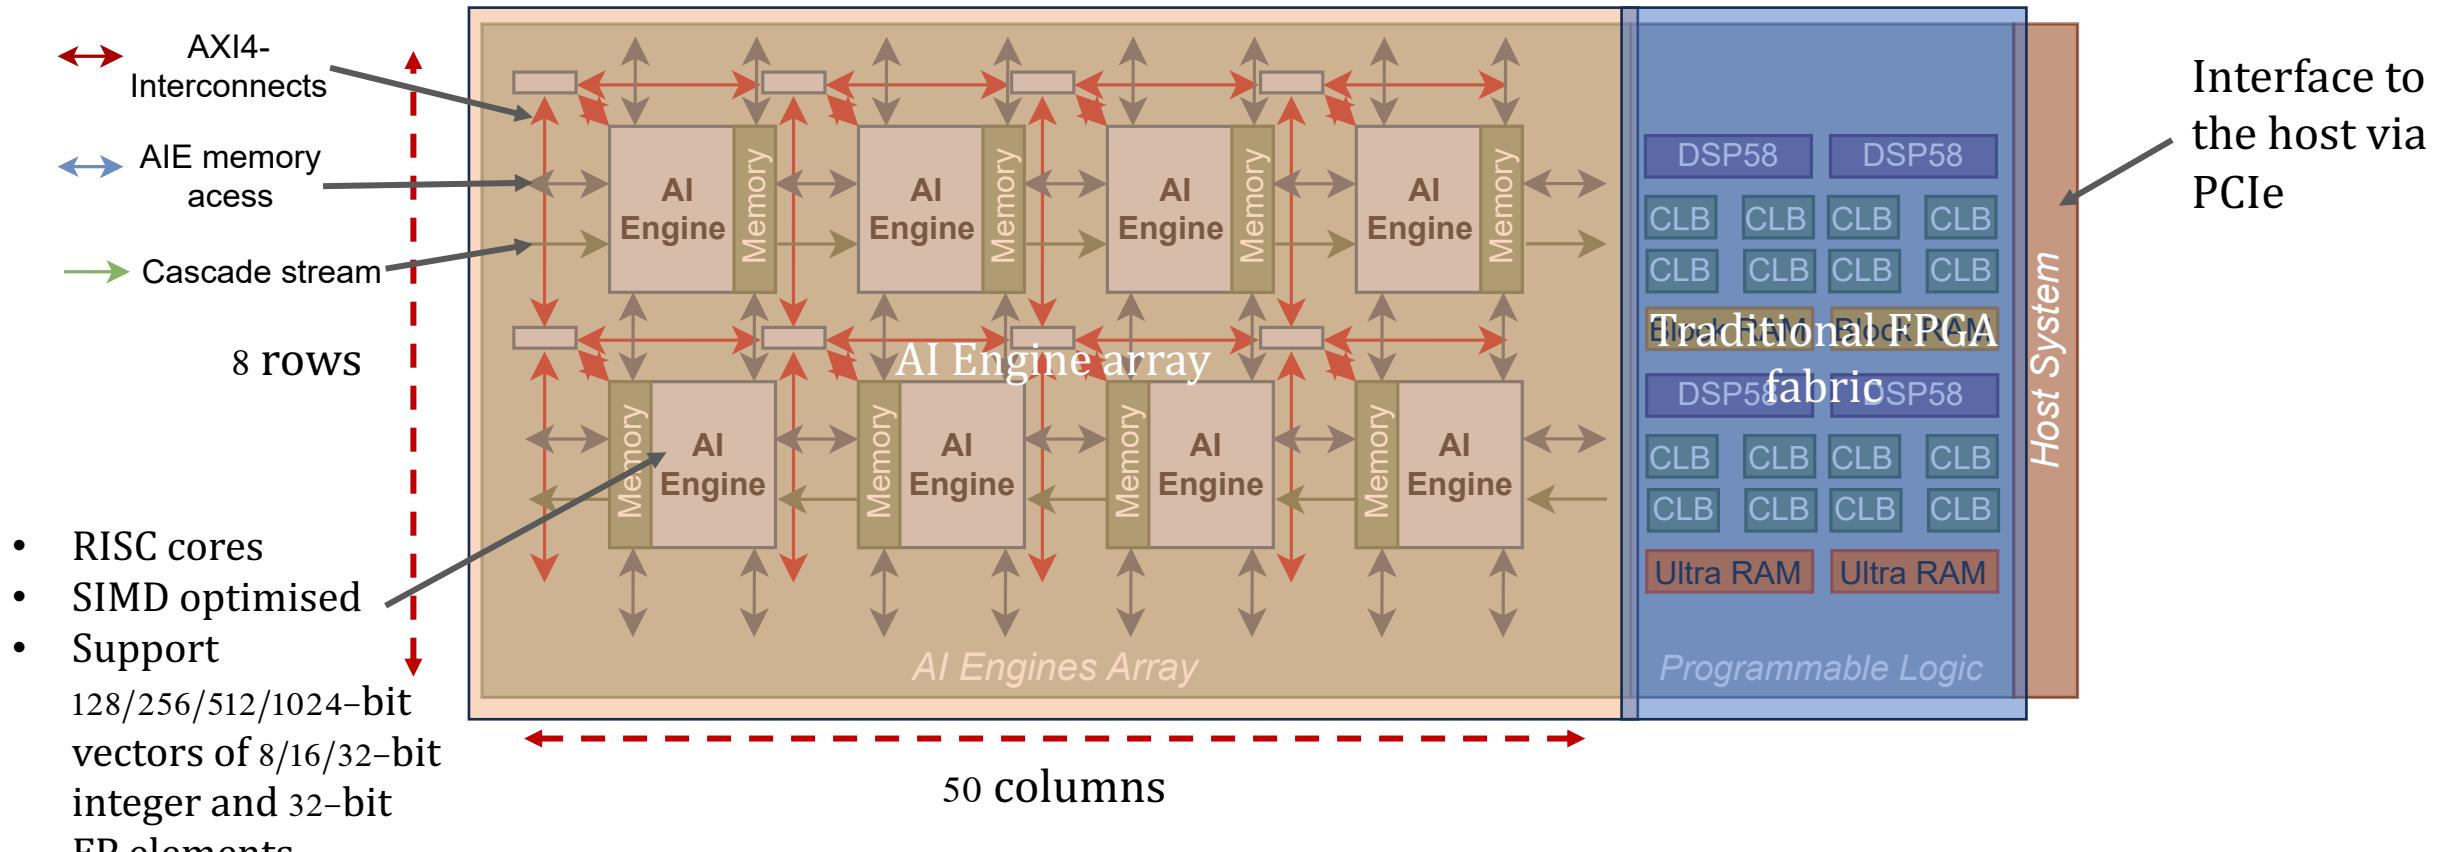


- PRNG implementation on Ambric AM2045 → An example mapping PRNG on a coarse-grained RISC-based architecture

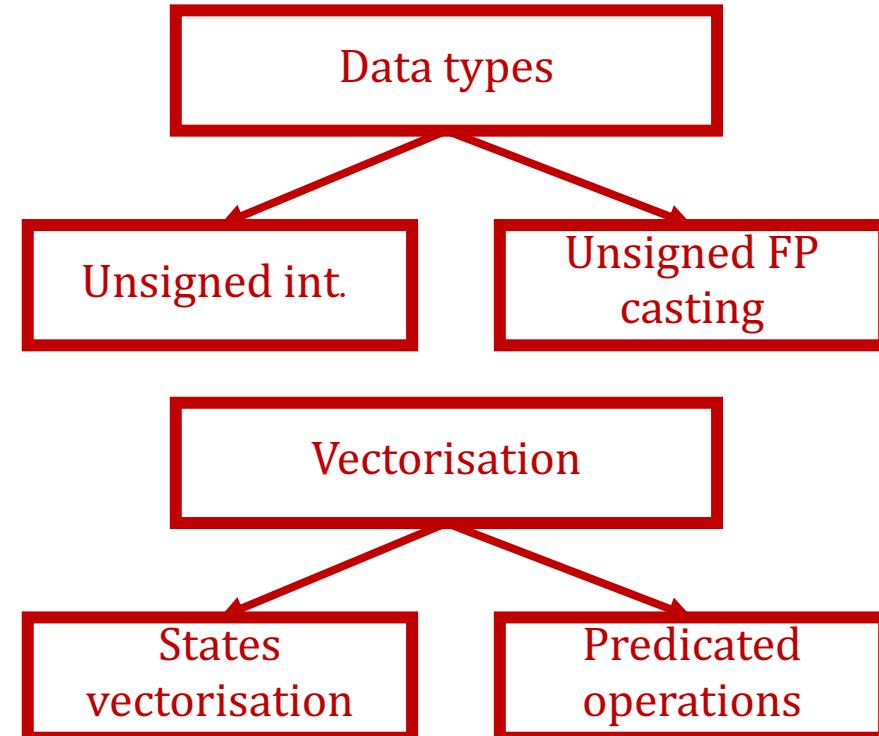
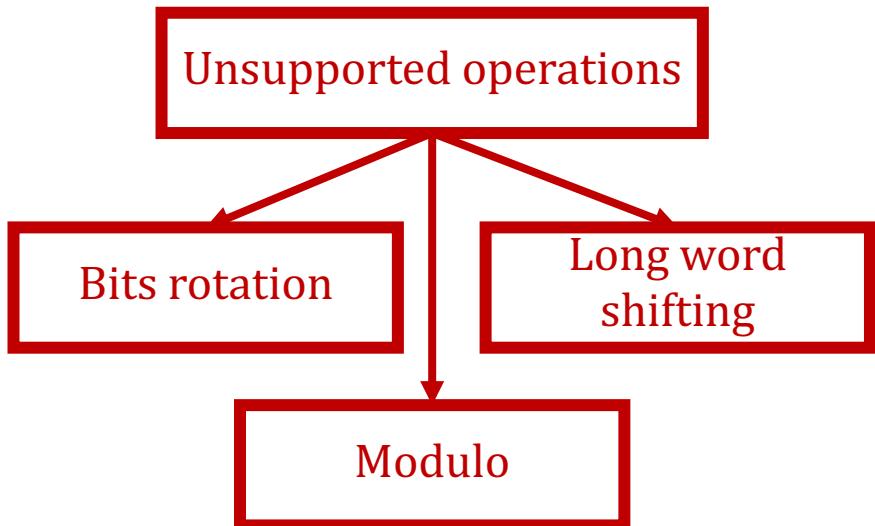


*AMD AI Engines architecture is also a RISC-based coarse-grained architecture with more capabilities!*

# AMD Versal SoC architecture



# Challenges in mapping PRNGs and Probability distribution approximations on AI Engines



# Main Contributions

- Propose an efficient implementation of the missed features.
- Propose the dataflow overlays of two execution models:
  - Co-processor model
  - Standalone accelerator model

# PRNG Execution Requirements

	<b>Data type</b>	<b>Required operations</b>	<b>Vectorisation</b>
<i>CUDA's by-default PRNG</i>	XORWOW	32-bit unsigned int.	
	XORSHIRO	64-bit unsigned int.	64-bit operations
<i>Vectorised version of MT</i>	SFMT	128-bit unsigned int.	128-bit operations Modulo operation 128-bit mask

# PRNG Execution Requirements, contd...

---

## Algorithm 1 XORWOW Random Number Generation

---

```

1: Input:  $s_0, \dots, s_4, \#RNs$ 
2: Output:  $RNs$ 
3: for  $i \leftarrow 0$  to  $\#RNs$  do
4:    $cnt \leftarrow cnt + 362437$ 
5:    $value \leftarrow s_4$ 
6:    $value \leftarrow value \oplus (value >> 2)$ 
7:    $value \leftarrow value \oplus (value << 1)$ 
8:    $value \leftarrow value \oplus (s_0 \oplus (s_0 << 4))$ 
9:    $value \leftarrow value + cnt$             $\triangleright$  Weyl sequence
10:   $RNs[i] \leftarrow value$ 
11:   $s_4 \leftarrow s_3, s_3 \leftarrow s_2, \dots, s_1 \leftarrow s_0$ 
12:   $s_0 \leftarrow value$ 
13: end for

```

---

## Algorithm 2 XOROSHIRO Random Number Generation

---

```

1: Input:  $s_0, s_1, \#RNs$ 
2: Output:  $RNs$ 
3: for  $i \leftarrow 0$  to  $\#RNs$  do
4:    $RNs[i] \leftarrow \text{rotate\_left}(s_0 + s_1, 17) + s_0$ 
5:    $s_1 \leftarrow s_0 \oplus s_1$ 
6:    $s_0 \leftarrow \text{rotate\_left}(s_0, 49) \oplus s_1 \oplus (s_1 << 21)$ 
7:    $s_1 \leftarrow \text{rotate\_left}(s_1, 37)$ 
8: end for

```

---

*unsigned integer =>  
Bitwise operations  
are insensitive to  
the sign-bit*

- *Rotation op.*
- *64-bit ops.*

---

## Algorithm 3 SFMT Random Number Generation

---

```

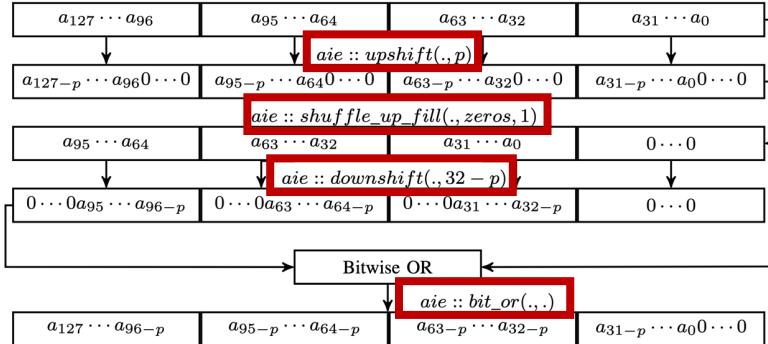
1: Macros: For MT19937 based algorithms,  $N = 156$  and
    $M = 122$ 
2: Input:  $s_0, s_1, \dots, s_N, \#RNs$ 
3: Output:  $RNs$ 
4:  $r_1 \leftarrow s_{N-2}, r_2 \leftarrow s_{N-1}$ 
5: for  $i \leftarrow 0$  to  $\#RNs$  do
6:    $v_0 \leftarrow s_i, v_M \leftarrow s_{(i+M)} \bmod N$ 
7:    $A \leftarrow (v_0 \overset{128}{<<} 8) \oplus v_0$ 
8:    $B \leftarrow (v_M \overset{32}{>>} 11) \& vMASK$ 
9:    $C \leftarrow r_1 \overset{128}{>>} 8$ 
10:   $D \leftarrow r_2 \overset{32}{<<} 18$ 
11:   $r_1 \leftarrow r_2$ 
12:   $r_2 \leftarrow A \oplus B \oplus C \oplus D$ 
13:   $RNs[i] \leftarrow r_2$ 
14:   $s_{N-1} \leftarrow r_2$ 
15: end for

```

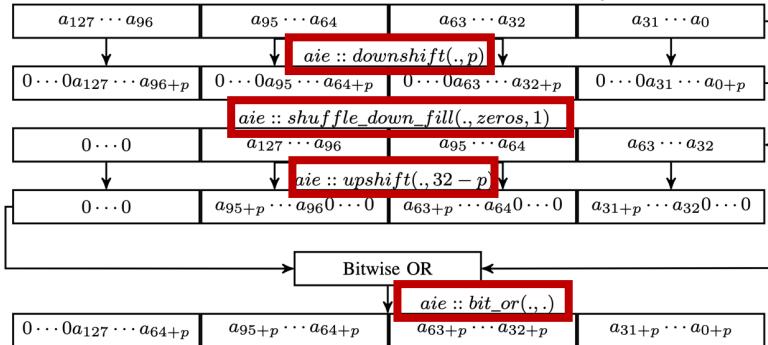
---

- *Long word shifting*
- *Modulo*
- *128-bit ops.*

# PRNG Execution Requirements, contd...

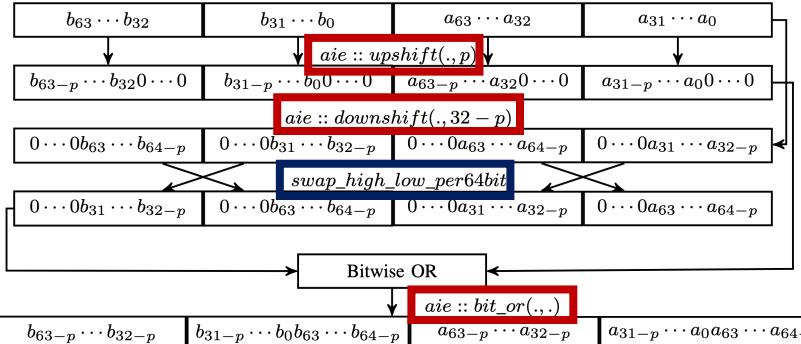


(a) Illustration of `sh1_128bit` function by  $p$  bits.

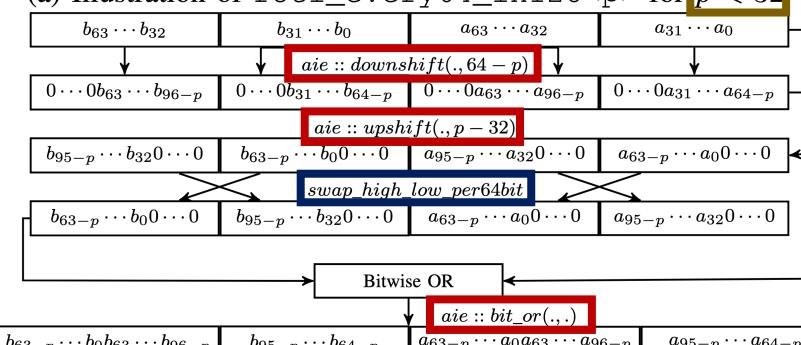


(b) Illustration of `shr_128bit` function by  $p$  bits.

Fig. 2: 128-bit shift operations.

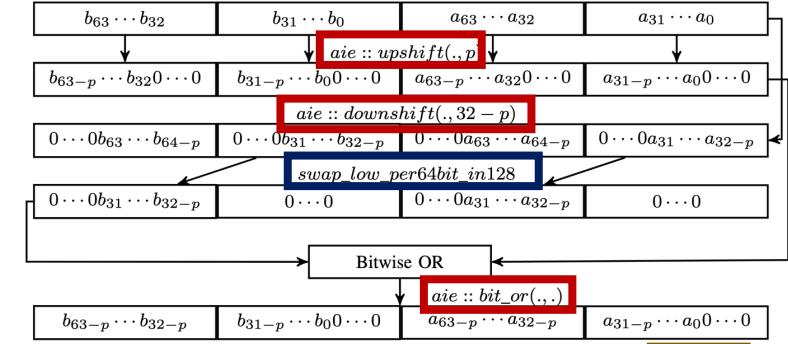


(a) Illustration of `rotl_every64_in128<p>` for  $p < 32$ .

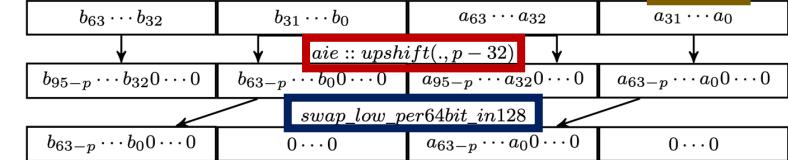


(b) Illustration of `rotl_every64_in128<p>` for  $p \geq 32$ .

Fig. 3: 2 × 64-bit rotate left operation.



(a) Illustration of `sh1_every64_in128<p>` for  $p < 32$ .



(b) Illustration of `sh1_every64_in128<p>` for  $p \geq 32$ .

Fig. 4: 2 × 64-bit left shift operation.

The same follows for the right shift

# Probability Distribution Approximations Requirements

Approximation of probability distribution



*Treats the integer as a 32-bit fixed-point number and converts it to 32-bit FP by specifying the position of the binary point.*

→ *By placing the binary point at the MSB, the numbers will be in [0,1)*

**BUT!!** The data types are considered signed!!

→ *Integers with MSB = 1 are misread as negative which results in negative FP.*

# Probability Distribution Approximations Requirements

Single-precision (32-bit) floating point representation



$$F = (-1)^{MSB} \times M \times 2^{E-128} \quad 1 \leq M < 2$$

32-bit Fixed-Point Number

*We want to extract the MSB as well*



# Probability Distribution Approximations Requirements

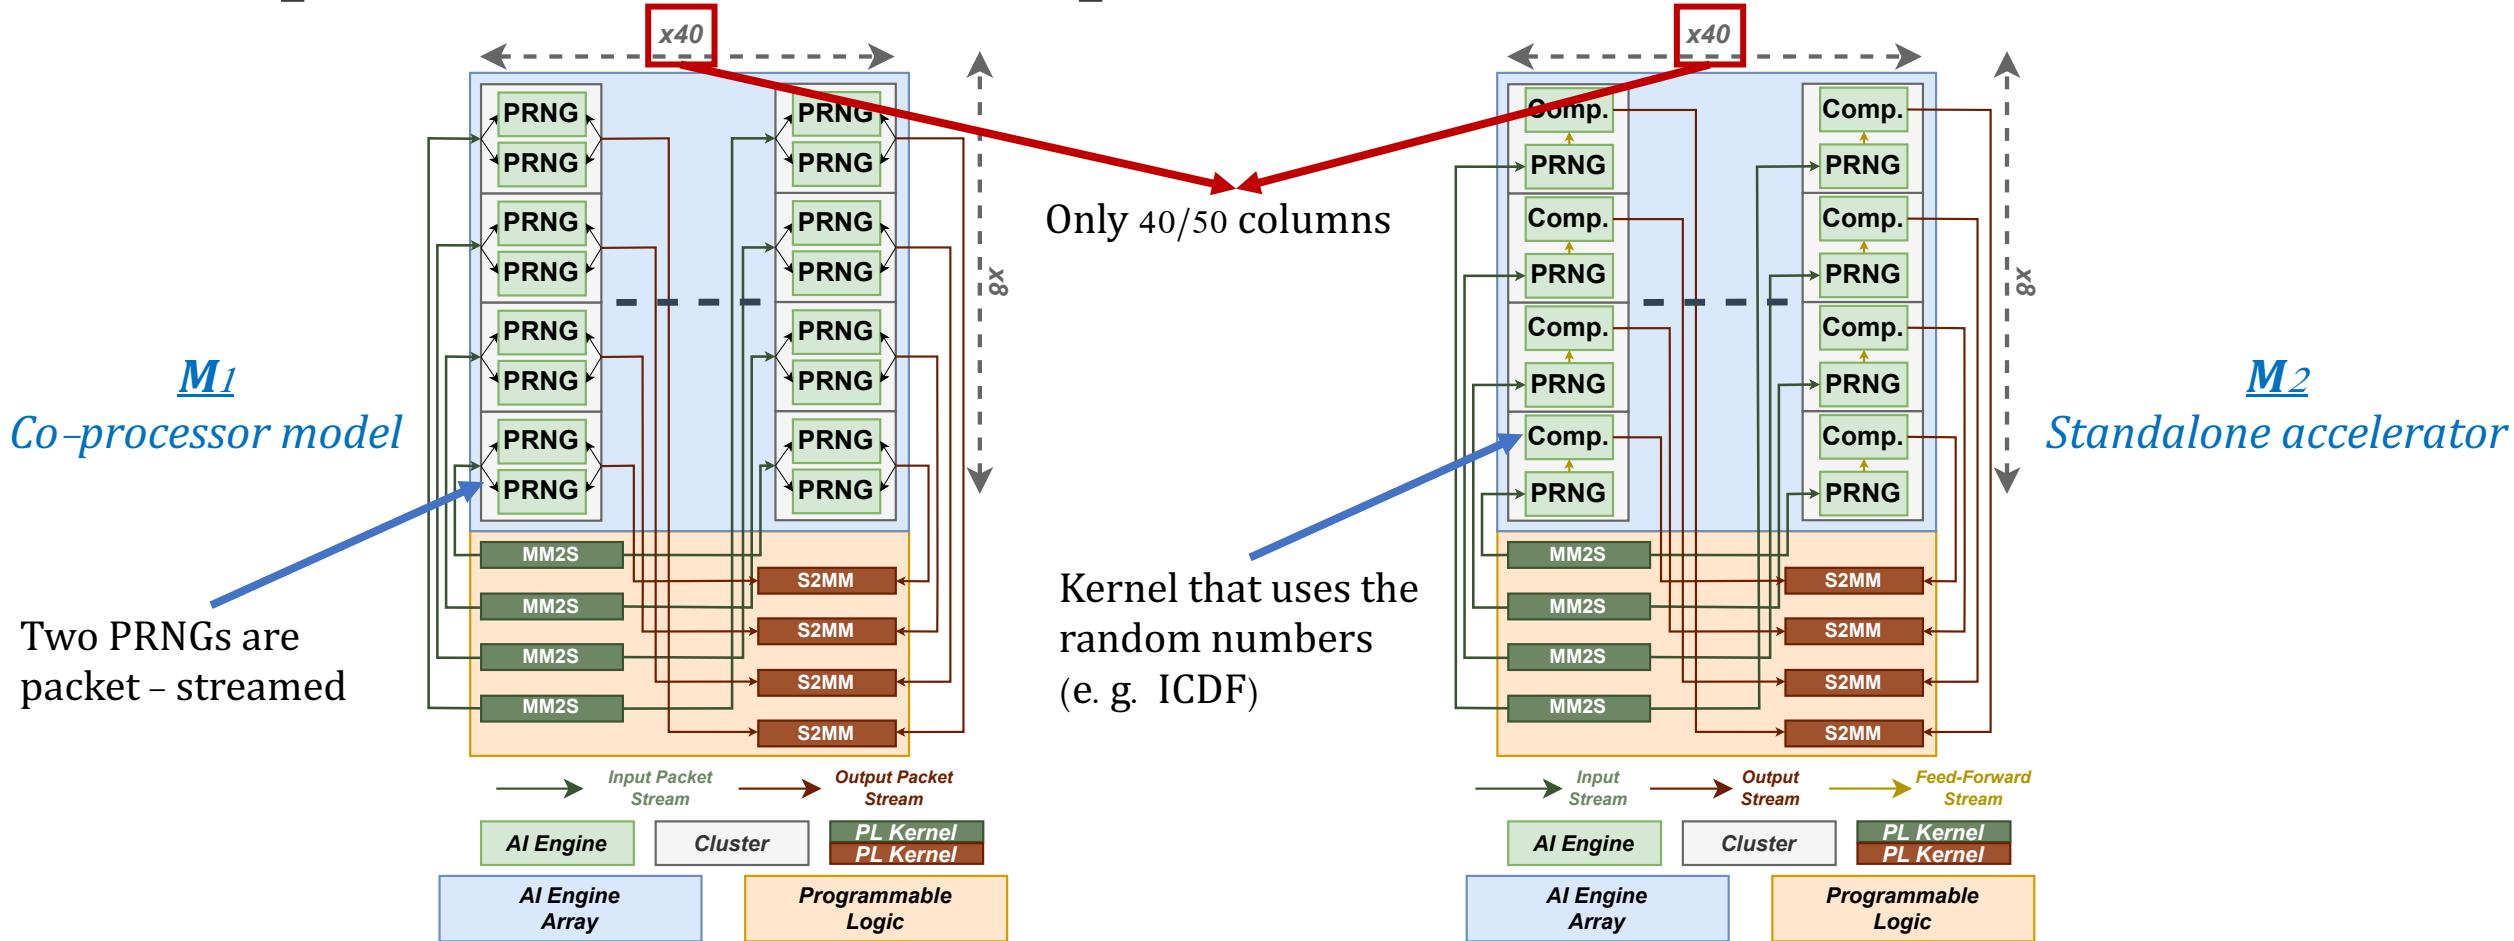


$$P(R(X)) = \frac{\left( \left( \left( \left( (A_1 R^2 + A_2) R^2 + A_3 \right) R^2 + A_4 \right) R^2 + A_5 \right) R^2 + A_6 \right) R}{\left( \left( \left( (B_1 R^2 + B_2) R^2 + B_3 \right) R^2 + B_4 \right) R^2 + B_5 \right) R^2 + 1}$$

- *Inner part of Acklam's approximation*
- *SIMD compliant*

where  $R(X) = X - 0.5$       *(coefficients in the paper)*

# Experimental Setup



# Experimental Results

PRNG	Model	160K	1.6M (10× ←)	16M (10× ←)	160M (10× ←)	M2 to M1
		RNs	RNs	RNs	RNs	
XORWOW	M1: Co-processor	2.858	17.477 (6.11×)	164.762 (9.42×)	1634.63 (9.92×)	3.15×
	M2: Dataflow	6.315	59.749 (9.46×)	593.921 (9.94×)	5934.01 (9.99×)	
SFMT	M1: Co-processor	4.831	30.494 (6.31×)	285.603 (9.36×)	2835.71 (9.92×)	1.86×
	M2: Dataflow	6.722	60.242 (8.96×)	594.355 (9.86×)	5934.5 (9.98×)	
XOROSHIRO	M1: Co-processor	7.637	70.596 (9.24×)	697.669 (9.88×)	6967.69 (9.98×)	0.96×
	M2: Dataflow	6.961	68.520 (9.84×)	683.628 (9.97×)	6833.75 (9.99×)	

- Scaling is sub-linear to linear
  - Overheads such as packet switching are well amortised (especially with XORWOW and SFMT at 1.6M)
  - XORWOW > SFMT > XOROSHIRO in speed in M1
  - XOROSHIRO ~ SFMT & XORWOW in speed in M2
- XORWOW and SFMT are suitable for M1 while XOROSHIRO is suitable for M2

# Conclusion

- PRNGs and probability distribution approximations can be part of dataflow processing on modern CGRAs/RDAs.
- AMD AI Engines is an example of such adaptive dataflow architectures.
- PRNGs and probability distribution approximations do not natively map to this architecture.
- Building the lacking features can be done using the available operations.
- Performance can be sub-linear to linear using these manipulations.
- Future work will focus on applying similar transformations on arithmetic operations to overcome the limitations of signed arithmetic with the Multiple-With-Carry PRNG family.

# Thanks!

**PRNGine**: Massively Parallel Pseudo-Random Number Generation and  
Probability Distribution Approximations on AMD AI Engines



[Github Repository](#)



[Reach out via email](#)



[Connect on LinkedIn](#)